

**Machine Learning in Passive Investment****Smart Beta Strategy**

Step 1. Rank all assets using exposure to each of the four metrics:  $\beta$

1. MOMENTUM
2. SIZE
  - book value (Book)
3. VALUE
  - trailing five-year average cash flow (Cash Flow)
  - trailing five-year average revenue (Revenue)
  - trailing five-year average gross sales (Sales)
  - trailing five-year average gross dividends (Dividends)
4. MACRO
  - total employment (Employment)
  - inflation (CPI or PCE)

Step 2. Select the 1,000 largest by each metric. Each of these 1,000 largest was included in the index at its relative metric weight  $\beta$ , in which way, we create the Fundamental index for that metric.

Step 3. The portfolio were rebalanced once every quarter, using only data available on the last trading day of the prior quarter.

**Smart Beta v.s. Active Alpha**

*Smart Beta* is often characterized as **passive investing** for the following reasons:

- It uses rule-based selection and weighting
- It rebalances at predetermined intervals
- **It does not attempt to make explicit forecasts of returns and risk for individual securities.**
- It requires additional active decisions to identify the specific factors to target, and to define the factors, the selection universe, the weighting method, and the rebalancing rules.
- All smart beta portfolios are concentrated in terms of the sources of excess return they seek to exploit, such as size, value, momentum, low volatility, or some combination of a few factors. This concentration leaves them susceptible to periods of poor returns to the chosen factors, which can lead to inconsistent performance.

*Active Alpha* strategy is active in several ways:

- It identifies numerous potentially profitable relationships between factors and prices.
- It forecasts returns to factors and individual securities, estimates associated risks, and revises forecasts based on the changing market environment.
- regularities in price response to a wide range of independent fundamental and behavioral factors. Portfolios are diversified across exposures to numerous factors.
- Risk models take on the same form as the alpha models. The pricing model interprets the equation by expectation, but the risk model interprets the equation by variance. The main challenge is that some factors may not be priced or rewarded unconditionally through time, but they do differentiate cross-sectional security returns. **These risk factors are mainly from three categories:**
  1. Macroeconomic factors (interest rates),
  2. Fundamental factors
  3. Statistical factors (principal component analysis).
- One can actively use the risk budget to construct portfolios instead of passively monitoring portfolio risk contribution (risk budgeting).

### **Unsupervised Learning and Smart Beta**

Smart beta strategy only selects one factor (metric) each time when the portfolio rebalances. In order to incorporate more than one factors (metrics, say momentum + value), one may use unsupervised learning methods to group stocks by their  $\beta$ s. In this part, we will discuss three main methods for clustering. These are:

- hierarchical clustering
- $k$ -mean
- model-based clustering

### **Hierarchical Clustering**

There are two types of hierarchical clustering methods: **agglomerative and divisive**.

- **Divisive methods** start with one large group and successively split the groups until there are  $n$  groups with one observation per group.
- **Agglomerative methods** are just the opposite; we start with  $n$  groups (one observation per group) and successively merge the two most similar groups until we are left with only one group. There are five commonly used methods for merging clusters in agglomerative clustering. These are single linkage, complete linkage, average linkage, centroid linkage and Ward's method.
  1. **The single linkage method** uses minimum distance, where the distance between clusters is defined as the distance between the closest pair of observations. Pairs consisting of one case from each group are used in the calculation. The first cluster is formed by merging the two groups with the shortest distance. Then the next smallest distance is found

between all of the clusters (keep in mind that an observation is also a cluster). The two clusters corresponding to the smallest distance are then merged.

2. The process for the **complete linkage method** is similar to single linkage, but the clustering criterion is different. The distance between groups is defined as the most distant pair of observations, with one coming from each group. The logic behind using this type of similarity criterion is that the maximum distance between observations in each cluster represents the smallest sphere that can enclose all of the objects in both clusters. Thus, the closest of these cluster pairs should be grouped together. The complete linkage method does not have the chaining problem that single linkage has.
3. **The average linkage method** for clustering starts out the same way as single and complete linkage. In this case, the cluster criterion is the average distance between all pairs, where one member of the pair comes from each cluster. Thus, we find all pairwise distances between observations in each cluster and take the average. This linkage method tends to combine clusters with small variances and to produce clusters with approximately equal variance.
4. **Centroid linkage** calculates the distance between two clusters as the distance between the centroids. The centroid of a cluster is defined as the  $d$ -dimensional sample mean for those observations that belong to the cluster. Whenever we merge clusters together or add an observation to a cluster, the centroid is recalculated.
5. The distance between two clusters using **Ward's linkage method** is defined as the incremental sum of the squares between two clusters. To merge clusters, the within-group sum-of-squares is minimized over all possible partitions obtained by combining two clusters. The within-group sum-of-squares is defined as the sum of the squared distances between all observations in a cluster and its centroid. This method tends to produce clusters with approximately the same number of observations in each one.

If the clustering is a valid one, then there should be a strong correlation between them. We can measure this using the **cophenetic correlation coefficient**. A cophenetic matrix is defined using the results of the linkage procedure. The  $ij$ -th entry of the cophenetic matrix is the fusion level at which the  $i$ -th and  $j$ -th objects appear together in the same cluster for the first time. The correlation coefficient between the distances and the corresponding cophenetic entries is the cophenetic correlation coefficient. Large values indicate that the linkage provides a reasonable clustering of the data.

## ***k*-Means Clustering**

The goal of  $k$ -means clustering is to partition the data into  $k$  groups such that the within-group sum-of-squares is minimized. One way this technique differs from hierarchical clustering is that we must specify the number of groups or clusters that we are looking for. We briefly describe two algorithms for obtaining clusters via  $k$ -means.

One of the basic algorithms for  $k$ -means clustering is a two step procedure. First, we assign observations to its closest group, usually using the Euclidean distance between the observation and the cluster centroid. The second step of the procedure is to calculate the new cluster centroid using the assigned objects. These steps are alternated until there are no changes in cluster membership or until the centroids do not change. This algorithm is sometimes referred to as HMEANS or the basic ISODATA method.

#### *PROCEDURE - HMEANS ALGORITHM*

1. Specify the number of clusters  $k$ .
2. Determine initial cluster centroids. These can be randomly chosen or the user can specify them.
3. Calculate the distance between each observation and each cluster centroid.
4. Assign every observation to the closest cluster.
5. Calculate the centroid (i.e., the  $d$ -dimensional mean) of every cluster using the observations that were just grouped there.
6. Repeat steps 3 through 5 until no more changes are made.

There are two problems with the HMEANS algorithm. The first one is that this method could lead to empty clusters, so users should be aware of this possibility. As the centroid is recalculated and observations are reassigned to groups, some clusters could become empty. The second issue concerns the optimality of the partitions. With  $k$ -means, we are searching for partitions where the within-group sum-of-squares is minimum. It can be shown that in some cases, the final  $k$ -means cluster assignment is not optimal, in the sense that moving a single point from one cluster to another may reduce the sum of squared errors. The following procedure helps address the second problem.

#### *PROCEDURE - K-MEANS*

1. Obtain a partition of  $k$  groups, possibly from the HMEANS algorithm.
2. Take each data point  $x_i$  and calculate the Euclidean distance between it and every cluster centroid.
3. Here  $x_i$  is in the  $r$ -th cluster,  $n_r$  is the number of points in the  $r$ -th cluster, and  $d_{ir}^2$  is the Euclidean distance between  $x_i$  and the centroid of cluster  $r$ . If there is a group  $s$  such that

$$\frac{n_r}{n_r - 1} d_{ir}^2 > \frac{n_s}{n_s + 1} d_{is}^2$$

then move  $x_i$  to cluster  $s$ .

4. If there are several clusters that satisfy the above inequality, then move the  $x_i$  to the group that has the smallest value for

$$\frac{n_s}{n_s + 1} d_{is}^2$$

5. Repeat steps 2 through 4 until no more changes are made.

There are also many algorithms for  $k$ -means clustering described in the literature.

## Model-Based Clustering

Model-based clustering combines several ideas:

- (1) **Finite Mixtures Models and EM Algorithm** is used to estimate the parameters of the finite mixture. Each component density of the finite mixture will represent a group or cluster.
- (2) **Model-based agglomerative clustering** is used to get a reasonable partition of the data. Each partition is then used to get initial starting values for the EM algorithm.
- (3) The **Bayes Information Criterion (BIC)** is used to choose the best grouping and number of clusters, given the data. The BIC is an approximation to Bayes factors.

### *Finite Mixture Models and the EM Algorithm*

#### Finite Mixture Model

We define the multivariate finite mixture model as the weighted sum of multivariate component densities:

$$f(x, p, \theta) = \sum_{k=1}^c p_k g_k(x, \theta_k)$$

The component density is denoted by  $g_k(x, \theta_k)$ , with parameters represented by  $\theta_k$ . The weights are given by  $p_k$ , with the constraint that they are nonnegative and sum to one. These weights are also called the mixing proportions or mixing coefficients.

The component density can be any bona fide probability density, but one of the most commonly used ones is the multivariate normal. This yields the following equation for a multivariate Gaussian finite mixture

$$f(x, p, \mu, \Sigma) = \sum_{k=1}^c p_k \phi(x, \mu_k, \Sigma_k)$$

where  $\phi$  represents a multivariate normal probability density function given by

$$\phi(x_i; \mu_k, \Sigma_k) = \frac{\exp\left\{-\frac{1}{2}(x - x_i)^T \Sigma_k^{-1}(x_i - \mu_k)\right\}}{(2\pi)^{d/2} \sqrt{|\Sigma_k|}}$$

$\mu_k$  is a  $d$ -dimensional vector of means, and  $\Sigma_k$  is a  $d \times d$  covariance matrix. To cluster the data we need to estimate the weights  $p_k$ , the  $d$ -dimension means for each component density and the covariance matrices. We have in total  $c \times d$  parameters for the means and  $c \times \frac{d(c+1)}{2}$  parameters for the component variance matrices.

#### EM Algorithm for Estimating Parameters

The problem of estimating the parameters in a finite mixture has been studied extensively in the literature. The book by Everitt and Hand (1981) provides an excellent overview of this topic and offers several methods for parameter estimation. The technique we present here is called the

Expectation-Maximization (EM) method. This is a general method for optimizing likelihood functions and is useful in situations where data might be missing or simpler optimization methods fail. The seminal paper on this topic is by Dempster, Laird and Rubin (1977), where they formalize the EM algorithm and establish its properties. Redner and Walker (1984) apply it to mixture densities. The EM methodology is now a standard tool for statisticians and is used in many applications.

In this section, we discuss the EM algorithm as it can be applied to estimating the parameters of a finite mixture of normal densities. To use the EM algorithm, we must have a value for the number of terms  $c$  in the mixture. This is usually obtained using prior knowledge of the application (the analyst expects a certain number of groups), using graphical exploratory data analysis (looking for clusters or other group structure) or using some other method of estimating the number of terms. The approach called adaptive mixtures [Priebe, 1994] offers a way to address the problem of determining the number of component densities to use in the finite mixture model. This approach is discussed later.

Besides the number of terms, we must also have an initial guess for the value of the component parameters. Once we have an initial estimate, we update the parameter estimates using the data and the equations given below. These are called the iterative EM update equations, and we provide the multivariate case as the most general one. The univariate case follows easily.

We wish to estimate

$$\theta = p_1, \dots, p_{c-1}, \mu_1, \dots, \mu_c, \Sigma_1, \dots, \Sigma_c$$

by maximizing the log-likelihood given by

$$\ln[L(\theta|x_1, \dots, x_n)] = \sum_{i=1}^n \ln \left[ \sum_{k=1}^c p_k \phi(x_i, \mu_k, \Sigma_k) \right]$$

Note that we don't have to estimate all  $c$  of the weight because

$$p_c = 1 - \sum_{i=1}^{c-1} p_i$$

we assume that the components exist in a fixed proportion in the mixture, given by the  $p_k$ .

The first step is to determine the posterior probabilities given by

$$\hat{\tau}_{ik}(x_i) = \frac{\hat{p}_k \phi(x_i; \hat{\mu}_k, \hat{\Sigma}_k)}{\hat{f}(x_i; \hat{p}, \hat{\mu}, \hat{\Sigma})}; \quad k = 1, \dots, c; \quad i = 1, \dots, n,$$

where  $\phi(x_i; \hat{\mu}_k, \hat{\Sigma}_k)$  is the multivariate normal density for the  $i$ -th term evaluated at  $x_j$ , and

$$\hat{f}(x_i; \hat{p}, \hat{\mu}, \hat{\Sigma}) = \sum_{j=1}^c \hat{p}_j \phi(x_i; \hat{\mu}_j, \hat{\Sigma}_j)$$

is the finite mixture estimate at point  $x_j$ .

The posterior probability tells us the likelihood that a point belongs to each of the separate component densities. We can use this estimated posterior probability to obtain a weighted update of the parameters for each component. This yields the iterative EM update equations for the mixing coefficients, the means and the covariance matrices. These are

$$\hat{p}_k = \frac{1}{n} \sum_{i=1}^n \hat{t}_{ik}$$

$$\hat{\mu}_k = \frac{1}{n} \sum_{i=1}^n \frac{\hat{t}_{ik} x_i}{\hat{p}_k}$$

$$\hat{\Sigma}_k = \frac{1}{n} \sum_{i=1}^n \frac{\hat{t}_{ik} (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)'}{\hat{p}_k}$$

Note that if  $d = 1$  then the update equation for the variance is

$$\hat{\sigma}_k^2 = \frac{1}{n} \sum_{i=1}^n \frac{\hat{t}_{ik} (x_i - \hat{\mu}_k)^2}{\hat{p}_k}$$

The EM algorithm is a two step process, consisting of an E-step and an M-step, as outline below. These two steps are repeated until the estimated value converge.

#### *FINITE MIXTURES - EM PROCEDURE*

1. Determine the number of terms or component densities  $c$  in the mixture.
2. Determine an initial guess at the component parameters. These are the mixing coefficients, means and covariance matrices for each normal density.
3. **(E-Step)** For each data point , calculate the posterior probability using

$$\hat{t}_{ik}(x_i) = \frac{\hat{p}_k \phi(x_i; \hat{\mu}_k, \hat{\Sigma}_k)}{\hat{f}(x_i; \hat{p}, \hat{\mu}, \hat{\Sigma})}; \quad k = 1, \dots, c; \quad i = 1, \dots, n,$$

4. **(M-Step)** Update the parameter estimates using the posterior probability from the E-step and equations

$$\begin{aligned} \hat{p}_k &= \frac{1}{n} \sum_{i=1}^n \hat{t}_{ik} \\ \hat{\mu}_k &= \frac{1}{n} \sum_{i=1}^n \frac{\hat{t}_{ik} x_i}{\hat{p}_k} \\ \hat{\sigma}_k^2 &= \frac{1}{n} \sum_{i=1}^n \frac{\hat{t}_{ik} (x_i - \hat{\mu}_k)^2}{\hat{p}_k} \end{aligned}$$

5. Repeat steps 3 through 4 until the estimates converge.

Typically, step 5 is implemented by continuing the iteration until the changes in the estimates at each iteration are less than some pre-set tolerance. Note that with the iterative EM algorithm, we need to use the entire data set to simultaneously update the parameter estimates. This imposes a high computational load when dealing with massive data sets.

**Model-based agglomerative clustering** uses the same general ideas as the agglomerate clustering we described earlier. In particular, all observations start out in their own singleton group, and two groups are merged at each step until we have just one group. It also provides a partition of the data for any given number of clusters between one and  $n$ . The difference between the two methods has to do with the criterion used to decide which two groups to merge at each step. With regular agglomerate clustering, the two closet groups (in terms of selected distance and type of linkage) are merged. In the case of model-based agglomerate clustering, we do not use a distance. Instead, we use the **classification likelihood** as our objective function, which is given by

$$L_{CL}(\theta_k, \gamma_i; x_i) = \prod_{i=1}^n f_{\gamma_i}(x_i; \theta_{\gamma_i})$$

where  $\gamma_i$  is a label indicating a classification for the  $i$ -th observation. we have  $\gamma_i = k$ , if  $x_i$  belongs to the  $k$ -th component. In the mixture approach, the number of observations in each component has a multinomial distribution with the sample size of  $n$  and probability parameters given by  $p_1, \dots, p_c$ . The exact form of classification likelihood will change depending on the model chosen for our finite mixture.

In the model based agglomerative clustering we seek to find partition that maximize (approximately) the classification likelihood. To accomplish this, the two clusters producing the largest increase in the classification likelihood are merged at each step. This process continues until all observations are in one group.

### **Bayesian Information Criterion**

We now turn our attention to the third major component in the model-based clustering, which is choosing the model that best fits our data according to some criterion.

The Bayesian Information Criterion is given by

$$BIC = 2 \log[L_M(X, \hat{\theta})] - m \log(n),$$

where  $L_M$  is the likelihood given that data, the model  $M$ , and the estimated parameters  $\hat{\theta}$ . The number of independent parameters to be estimated for model  $M$  is given by  $m$ .